

# Large Scale Deep Learning and Self-Supervision in Vision and NLP

Thomas Breuel

NVIDIA

# INTRODUCTION

# Topics

- using large amounts of unlabeled training data
- self-supervised training
- LSTM, transformers, ViT
- large language models
- diffusion models

Application: OCR (text recognition, layout recognition, information extraction)

# Different Kinds of Learning

- supervised learning
  - inputs and outputs are given
- unsupervised learning
  - only inputs are given
- semi-supervised learning
  - combine supervised and unsupervised data
- self-supervised learning
  - algorithm derives a supervised problem from unsupervised data

# Different Kinds of Learning

- transfer learning
  - use a model trained on one problem to solve another problem
- active learning
  - the learning algorithm requests transcribed data from an oracle
- metric learning
  - learn a distance measure that helps with clustering / classification
- representation learning
  - transform input vectors into another space that, removes noise, expresses invariances, places "similar" samples closer together

# Different Kinds of Models

- discriminative vs generative
- linear models
- non-linear models
  - parametric probabilistic
  - support vector machines
  - deep learning
    - convolutional
    - recurrent
    - transformer

# Overview

Motivation:

- OCR example

Classical Techniques:

- statistical basis of machine learning (mostly review)
- linear methods (mostly review?)

Deep Learning:

- language modeling and sequence learning
- self-supervised learning for images
- generative modeling (VAE, flows, GANs)

# Background

I'm assuming you have

- done some supervised deep learning, e.g., trained MNIST, ImageNet, etc.
- understand the basics of SGD, loss functions, etc.



# Focus

We will focus on the unsupervised / self-supervised aspects of methods and papers we will be discussing, discussing other aspects of papers only as needed.