

Almahairi, Amjad, Sai Rajeshwar, Alessandro Sordoni, Philip Bachman, and Aaron Courville. "Augmented CycleGAN: Learning Many-to-Many Mappings from Unpaired Data." In *Proceedings of the 35th International Conference on Machine Learning*, 195–204. PMLR, 2018. <https://proceedings.mlr.press/v80/almahairi18a.html>.

Arjovsky, Martin, Soumith Chintala, and Léon Bottou. "Wasserstein GAN." arXiv, December 6, 2017. <https://doi.org/10.48550/arXiv.1701.07875>.

Bao, Hangbo, Li Dong, Songhao Piao, and Furu Wei. "BEiT: BERT Pre-Training of Image Transformers." arXiv, September 3, 2022. <https://doi.org/10.48550/arXiv.2106.08254>.

Bell, Anthony J., and Terrence J. Sejnowski. "An Information-Maximization Approach to Blind Separation and Blind Deconvolution." *Neural Computation* 7 (1995): 1129–59.

BELL, ANTHONY J., and TERRENCE J. SEJNOWSKI. "The 'Independent Components' of Natural Scenes Are Edge Filters." *Vision Research* 37, no. 23 (December 1997): 3327–38.

Brock, Andrew, Jeff Donahue, and Karen Simonyan. "Large Scale GAN Training for High Fidelity Natural Image Synthesis." arXiv, February 25, 2019. <https://doi.org/10.48550/arXiv.1809.11096>.

Bromley, Jane, Isabelle Guyon, Yann LeCun, Eduard Säckinger, and Roopak Shah. "Signature Verification Using a 'Siamese' Time Delay Neural Network," n.d., 8.

Brown, Tom B., Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, et al. "Language Models Are Few-Shot Learners." arXiv, July 22, 2020. <https://doi.org/10.48550/arXiv.2005.14165>.

———. "Language Models Are Few-Shot Learners." arXiv, July 22, 2020. <https://doi.org/10.48550/arXiv.2005.14165>.

Caron, Mathilde, Hugo Touvron, Ishan Misra, Hervé Jégou, Julien Mairal, Piotr Bojanowski, and Armand Joulin. "Emerging Properties in Self-Supervised Vision Transformers." arXiv, May 24, 2021. <https://doi.org/10.48550/arXiv.2104.14294>.

Chen, Ting, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. "A Simple Framework for Contrastive Learning of Visual Representations." arXiv, June 30, 2020. <https://doi.org/10.48550/arXiv.2002.05709>.

Chen, Ting, Simon Kornblith, Kevin Swersky, Mohammad Norouzi, and Geoffrey Hinton. "Big Self-Supervised Models Are Strong Semi-Supervised Learners." arXiv, October 25, 2020. <https://doi.org/10.48550/arXiv.2006.10029>.

Chen, Xi, Yan Duan, Rein Houthooft, John Schulman, Ilya Sutskever, and Pieter Abbeel. "InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets." arXiv, June 11, 2016. <https://doi.org/10.48550/arXiv.1606.03657>.

Creswell, Antonia, Tom White, Vincent Dumoulin, Kai Arulkumaran, Biswa Sengupta, and Anil A. Bharath. "Generative Adversarial Networks: An Overview." *IEEE Signal Processing Magazine* 35, no. 1 (January 2018): 53–65. <https://doi.org/10.1109/MSP.2017.2765202>.

Croitoru, Florinel-Alin, Vlad Hondru, Radu Tudor Ionescu, and Mubarak Shah. "Diffusion Models in Vision: A Survey." arXiv, September 10, 2022. <https://doi.org/10.48550/arXiv.2209.04747>.

Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. "BERT: Pre-Training of Deep Bidirectional Transformers for Language Understanding." arXiv, May 24, 2019. <https://doi.org/10.48550/arXiv.1810.04805>.

Dinh, Laurent, David Krueger, and Yoshua Bengio. "NICE: Non-Linear Independent Components Estimation." arXiv, April 10, 2015. <https://doi.org/10.48550/arXiv.1410.8516>.

Dinh, Laurent, Jascha Sohl-Dickstein, and Samy Bengio. "Density Estimation Using Real NVP." arXiv, February 27, 2017. <https://doi.org/10.48550/arXiv.1605.08803>.

Doersch, Carl, Abhinav Gupta, and Alexei A. Efros. "Unsupervised Visual Representation Learning by Context Prediction." arXiv, January 16, 2016. <https://doi.org/10.48550/arXiv.1505.05192>.

Dosovitskiy, Alexey, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, et al. "An Image Is Worth 16x16 Words: Transformers for Image Recognition at Scale." arXiv, June 3, 2021. <https://doi.org/10.48550/arXiv.2010.11929>.

Du, Yilun, and Igor Mordatch. "Implicit Generation and Generalization in Energy-Based Models." arXiv, June 29, 2020. <https://doi.org/10.48550/arXiv.1903.08689>.

Goldberg, Yoav, and Omer Levy. "Word2vec Explained: Deriving Mikolov et al.'s Negative-Sampling Word-Embedding Method." arXiv, February 15, 2014. <https://doi.org/10.48550/arXiv.1402.3722>.

Goodfellow, Ian J., Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. "Generative Adversarial Networks." arXiv, June 10, 2014. <https://doi.org/10.48550/arXiv.1406.2661>.

Graves, Alex. "Generating Sequences With Recurrent Neural Networks." arXiv, June 5, 2014. <http://arxiv.org/abs/1308.0850>.

Grill, Jean-Bastien, Florian Strub, Florent Altché, Corentin Tallec, Pierre H. Richemond, Elena Buchatskaya, Carl Doersch, et al. "Bootstrap Your Own Latent: A New Approach to Self-Supervised Learning." arXiv, September 10, 2020. <https://doi.org/10.48550/arXiv.2006.07733>.

He, Kaiming, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. "Masked Autoencoders Are Scalable Vision Learners." arXiv, December 19, 2021. <https://doi.org/10.48550/arXiv.2111.06377>.

OpenAI. "Implicit Generation and Generalization Methods for Energy-Based Models," March 21, 2019. <https://openai.com/blog/energy-based-models/>.

Jang, Eric, Shixiang Gu, and Ben Poole. "Categorical Reparameterization with Gumbel-Softmax." arXiv, August 5, 2017. <https://doi.org/10.48550/arXiv.1611.01144>.

Kingma, Diederik P., and Max Welling. "An Introduction to Variational Autoencoders." *Foundations and Trends® in Machine Learning* 12, no. 4 (2019): 307–92. <https://doi.org/10.1561/22000000056>.

———. "An Introduction to Variational Autoencoders." *Foundations and Trends® in Machine Learning* 12, no. 4 (2019): 307–92. <https://doi.org/10.1561/22000000056>.

———. "Auto-Encoding Variational Bayes." arXiv, May 1, 2014. <https://doi.org/10.48550/arXiv.1312.6114>.

Koch, Gregory, Richard Zemel, and Ruslan Salakhutdinov. "Siamese Neural Networks for One-Shot Image Recognition,"

- n.d., 8.
- LeCun, Yann. "Learning Invariant Feature Hierarchies." In *Computer Vision – ECCV 2012. Workshops and Demonstrations*, edited by Andrea Fusiello, Vittorio Murino, and Rita Cucchiara, 7583:496–505. Lecture Notes in Computer Science. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012. https://doi.org/10.1007/978-3-642-33863-2_51.
- Makansi, Osama, Eddy Ilg, Özgün Cicek, and Thomas Brox. "Overcoming Limitations of Mixture Density Networks: A Sampling and Fitting Framework for Multimodal Future Prediction." In *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 7137–46, 2019. <https://doi.org/10.1109/CVPR.2019.00731>.
- McCann, Bryan, James Bradbury, Caiming Xiong, and Richard Socher. "Learned in Translation: Contextualized Word Vectors." arXiv, June 20, 2018. <https://doi.org/10.48550/arXiv.1708.00107>.
- Mikolov, Tomas, Kai Chen, Greg Corrado, and Jeffrey Dean. "Efficient Estimation of Word Representations in Vector Space." arXiv, September 6, 2013. <https://doi.org/10.48550/arXiv.1301.3781>.
- Oord, Aaron van den, Oriol Vinyals, and Koray Kavukcuoglu. "Neural Discrete Representation Learning." arXiv, May 30, 2018. <http://arxiv.org/abs/1711.00937>.
- Papamakarios, George, Eric Nalisnick, Danilo Jimenez Rezende, Shakir Mohamed, and Balaji Lakshminarayanan. "Normalizing Flows for Probabilistic Modeling and Inference." arXiv, April 8, 2021. <https://doi.org/10.48550/arXiv.1912.02762>.
- Park, Taesung, Alexei A. Efros, Richard Zhang, and Jun-Yan Zhu. "Contrastive Learning for Unpaired Image-to-Image Translation." arXiv, August 20, 2020. <https://doi.org/10.48550/arXiv.2007.15651>.
- Pathak, Deepak, Philipp Krahenbuhl, Jeff Donahue, Trevor Darrell, and Alexei A. Efros. "Context Encoders: Feature Learning by Inpainting." arXiv, November 21, 2016. <https://doi.org/10.48550/arXiv.1604.07379>.
- Pennington, Jeffrey, Richard Socher, and Christopher Manning. "Glove: Global Vectors for Word Representation." In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 1532–43. Doha, Qatar: Association for Computational Linguistics, 2014. <https://doi.org/10.3115/v1/D14-1162>.
- Pentland, Alex. "E i g e d c e s for Recognition," n.d., 16.
- Peters, Matthew E., Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. "Deep Contextualized Word Representations." arXiv, March 22, 2018. <http://arxiv.org/abs/1802.05365>.
- Principe, Jose C., Dongxin Xu, and John W. Fisher Iii. "Information-Theoretic Learning," 1999.
- Radford, Alec, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. "Improving Language Understanding by Generative Pre-Training," n.d., 12.
- Razavi, Ali, Aaron van den Oord, and Oriol Vinyals. "Generating Diverse High-Fidelity Images with VQ-VAE-2." arXiv, June 2, 2019. <https://doi.org/10.48550/arXiv.1906.00446>.
- Rocca, Joseph. "Understanding Variational Autoencoders (VAEs)." Medium, March 21, 2021. <https://towardsdatascience.com/understanding-variational-autoencoders-vaes-f70510919f73>.
- Shree, Priya. "The Journey of Open AI GPT Models." *Walmart Global Tech Blog* (blog), November 10, 2020. <https://medium.com/walmartglobaltech/the-journey-of-open-ai-gpt-models-32d95b7b7fb2>.
- Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. "Attention Is All You Need." arXiv, December 5, 2017. <https://doi.org/10.48550/arXiv.1706.03762>.
- Vincent, Pascal, Hugo Larochelle, Isabelle Lajoie, Yoshua Bengio, and Pierre-Antoine Manzagol. "Stacked Denoising Autoencoders: Learning Useful Representations in a Deep Network with a Local Denoising Criterion," n.d., 38.
- Zhang, Han, Ian Goodfellow, Dimitris Metaxas, and Augustus Odena. "Self-Attention Generative Adversarial Networks." arXiv, June 14, 2019. <https://doi.org/10.48550/arXiv.1805.08318>.
- Zhang, Richard, Phillip Isola, and Alexei A. Efros. "Split-Brain Autoencoders: Unsupervised Learning by Cross-Channel Prediction." arXiv, April 20, 2017. <https://doi.org/10.48550/arXiv.1611.09842>.
- Zhu, Jun-Yan, Taesung Park, Phillip Isola, and Alexei A. Efros. "Unpaired Image-to-Image Translation Using Cycle-Consistent Adversarial Networks." arXiv, August 24, 2020. <https://doi.org/10.48550/arXiv.1703.10593>.